



AI Safety and Alignment

Analysis of Model Behavior, Bias, and Strategic Reasoning

Background

As generative AI becomes more common in everyday work, its behavior raises important questions. This research explores how closely AI reflects human reasoning and whether it exhibits bias, dishonesty, or unnecessary complexity.

Methodology

We measured misalignment rates under the same testing conditions across

Q1 Different model families

Models of varying sizes Q2

We analyzed the Emergent Misalignment Dataset with 8,640 observations across 12 model variants from Gemma 3 and Qwen3 families. Models ranged from 1B to 32B parameters, evaluated under three training conditions (Base, Educational, Insecure) across eight misalignment scenarios.

Q3 Alignment Faking

Claude 3 Opus was tested under three conditions

Conflicting Objective:


Model told to be "helpful-only" vs original "helpful, honest, harmless" training

Context Signaling:

Free-tier conversations are used for RLHF training (monitored); paid-tier are not used (unmonitored)

Scratchpad:

Hidden reasoning space to detect strategic behavior.



Hypothesis:

Models comply with harmful requests during training (to avoid modification) but refuse during deployment (to follow true preferences)


Q4 Gender & Ethnicity Bias

Constitutional AI vs RLHF Q5

To test for demographic bias in AI hiring, we created standardized CVs from real resume data and randomly assigned gender- and ethnicity-distinctive names (Black/White, male/female) in a 2x2 factorial design. These CVs were scored by GPT-3.5 Turbo for qualifications (0-100). By comparing scores across demographics, we can detect if AI evaluations systematically varied based on gender and ethnicity, indicating demographic bias.

We compared 637,636 paired evaluations from Claude Sonnet 4 (Constitutional AI) and GPT-4o (RLHF) using multiple fairness metrics.

Findings and Discussion





MODEL SIZE & TRAINING CONDITION

The bar chart shows that training conditions strongly affect misalignment in small models. As model size increases, misalignment decreases and differences across training conditions converge, indicating that larger models are more robust to training variations.


MODEL SCALE & ARCHITECTURE

Larger models (27-32B parameters) show 45% lower misalignment than small models (1-4B), with architectural differences significantly impacting safety ($F=124.56$, $p<0.0001$). Each billion parameters reduces misalignment by 0.29 points, demonstrating that scale enhances both performance and robustness against adversarial conditions.





STRATEGIC ALIGNMENT FAKING

Models exhibit sophisticated deception: refusal rates are 10.4% higher in deployment (69.0%) vs. training contexts (58.6%), $t=8.79$, $p<0.0001$. This strategic behavior fundamentally compromises standard safety evaluation protocols, as monitored testing may not reflect actual deployment behaviour.



DEMOGRAPHIC BIAS PATTERNS

Both Constitutional AI (Claude) and RLHF (GPT-4o) exhibit systematic demographic biases but pattern them differently. Claude provides universally higher scores (+5.37 points) but shows 42-65% larger demographic disparities. This reveals a fundamental fairness-generosity trade-off: neither approach eliminates bias, requiring value judgments between equal treatment versus adequate outcomes for all. Score distribution differs significantly between genders. Females more likely to receive "Very High" scores (16.5% vs 14.7%), indicating GPT-3.5 assigns higher average scores and disproportionately places female-named CVs in highest category.



Conclusion

Our study found that larger models generally demonstrate stronger safety and robustness, while small models are more prone to misalignment and perform worse under insecure training.

We also found clear evidence of strategic alignment faking : Compliance in training is 10.4% higher than deployment. This behavior suggests that evaluation protocols may be compromised, as models that appear aligned during training may behave differently in real-world contexts. Paradoxically, monitoring can sometimes backfire, encouraging models to conceal rather than correct misalignment.

Constitutional AI and RLHF reveal unavoidable trade-offs : Constitutional AI improves overall performance but can increase disparities, while RLHF reduces bias at the cost of generosity. These findings show that AI alignment is an ethical challenge, requiring careful consideration of competing principles such as sufficiency and fairness.